


ECE3073

Real Time Systems Utilisation Statistics and Why do we need them?

Clive Maynard © 2023



Acknowledgement of Country



We wish to acknowledge the people of the Kulin Nations, on whose land Monash University operates.

We pay our respects to their Elders, past and present.

WARNING

COMMONWEALTH OF AUSTRALIA Copyright Regulations 1969

This material has been reproduced and communicated to you by or on behalf of Monash University pursuant to Part VB of the Copyright Act 1968 (the Act).

The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice

Starting Multitasking

```
void main (void)
{
    INT8U err;
    OSInit();

    /* ----- Application Initialization ----- */
    OSMutexCreate(9, &err);
    OSTaskCreate(TaskPrio10, (void *)0, &TaskPrio10Stk[999], 10);
    OSTaskCreate(TaskPrio15, (void *)0, &TaskPrio15Stk[999], 15);
    OSTaskCreate(TaskPrio20, (void *)0, &TaskPrio20Stk[999], 20);
    /* ----- Application Initialization ----- */

    OSStart();
}
```

Starting Multitasking(2)

```
void main(void)
{
    OsInit()
    OSTaskCreateExt(TaskStart, ...
    OSStart();
}
```

A related context for systems.

There are three working scenarios
to consider:

- 1) Need to know
- 2) Want to know
- 3) Access to knowledge

Utilisation

- The study of, and application, of scheduling algorithms requires knowledge of how much work the CPU does.
- The measure is “utilisation”....what percentage of the CPU power is used.
- How do we know the total utilisation and the utilisation of individual tasks?

Utilisation

- We can calculate it from the code and simulation.
- We can measure it directly.
- We can obtain it indirectly during execution by collecting statistical information about how much free time the CPU has available.

Calculating Utilisation

- Run your task codes through a simulator and from the total execution cycles you have a measure of each task utilisation.
- Summing this for all tasks and knowing system clock speed you can obtain the total utilisation and, therefore, the CPU load.

Measuring Utilisation

- With our tasks running on a target processor we can measure the execution time of a task (just as in the lab tests)
- For periodic tasks we know how often they need to run so the task utilisation is simply $(\text{Execution time}) / \text{Period}$
- For aperiodic tasks we need to know the worst case (highest) execution rate for the task

Statistical Evaluation

- How do we know how much spare time the CPU has?
- Collect data on how much time it is not doing anything useful.....the idle time
- The problem.....processor clock speed and timers compared with “real time”.
- How do we determine this in a processor independent fashion?

Timer Interrupt


- **Why does an RTOS need a timer interrupt?**
 - Sleeping: zzzzzzz....
 - Waking: !!!!!!!!!!!
 - Timeouts: whooops, give up
 - Keeping track of time: tick tock tick tock

Using the timer interrupt

- **OSTimeDlyHMSM(h,m,s,ms)**
 - Delay task for specified time
- **OSTimeDly(n)**
 - Delay task for specified number of ticks (timer interrupts)

uC OSTimeDly(n)

- Actual time delay may vary:



Collecting Statistics

- In an RTOS designed to be portable across many processors this is not an easy job.
- The designer of uCOSII (Jean Labrosse) solved it in a rather elegant way but to use it he needed a very particular approach.
- He should be proud of it and the next slides go through the process.
- To use it you just need to ensure you use the required approach.

uC OS- II Idle Task

- The **Idle task runs when no other task is ready**
- The **Idle task is the lowest priority task.**
- **Allows CPU usage to be measured using the counter OSIdleCtr:**

```
void OS_TaskIdle( void *pdata){  
    for (;;) {  
        OS_ENTER_CRITICAL();  
        OSIdleCtr++;  
        OS_EXIT_CRITICAL();  
        OSTaskIdleHook();  
    }  
}
```


Statistics Task

- The task **OS_TaskStat()** runs every second
- Calculates integer % CPU usage stored in **OSCPUUsage**
- **OSCPUUsage** is a global variable, so accessible from any task.

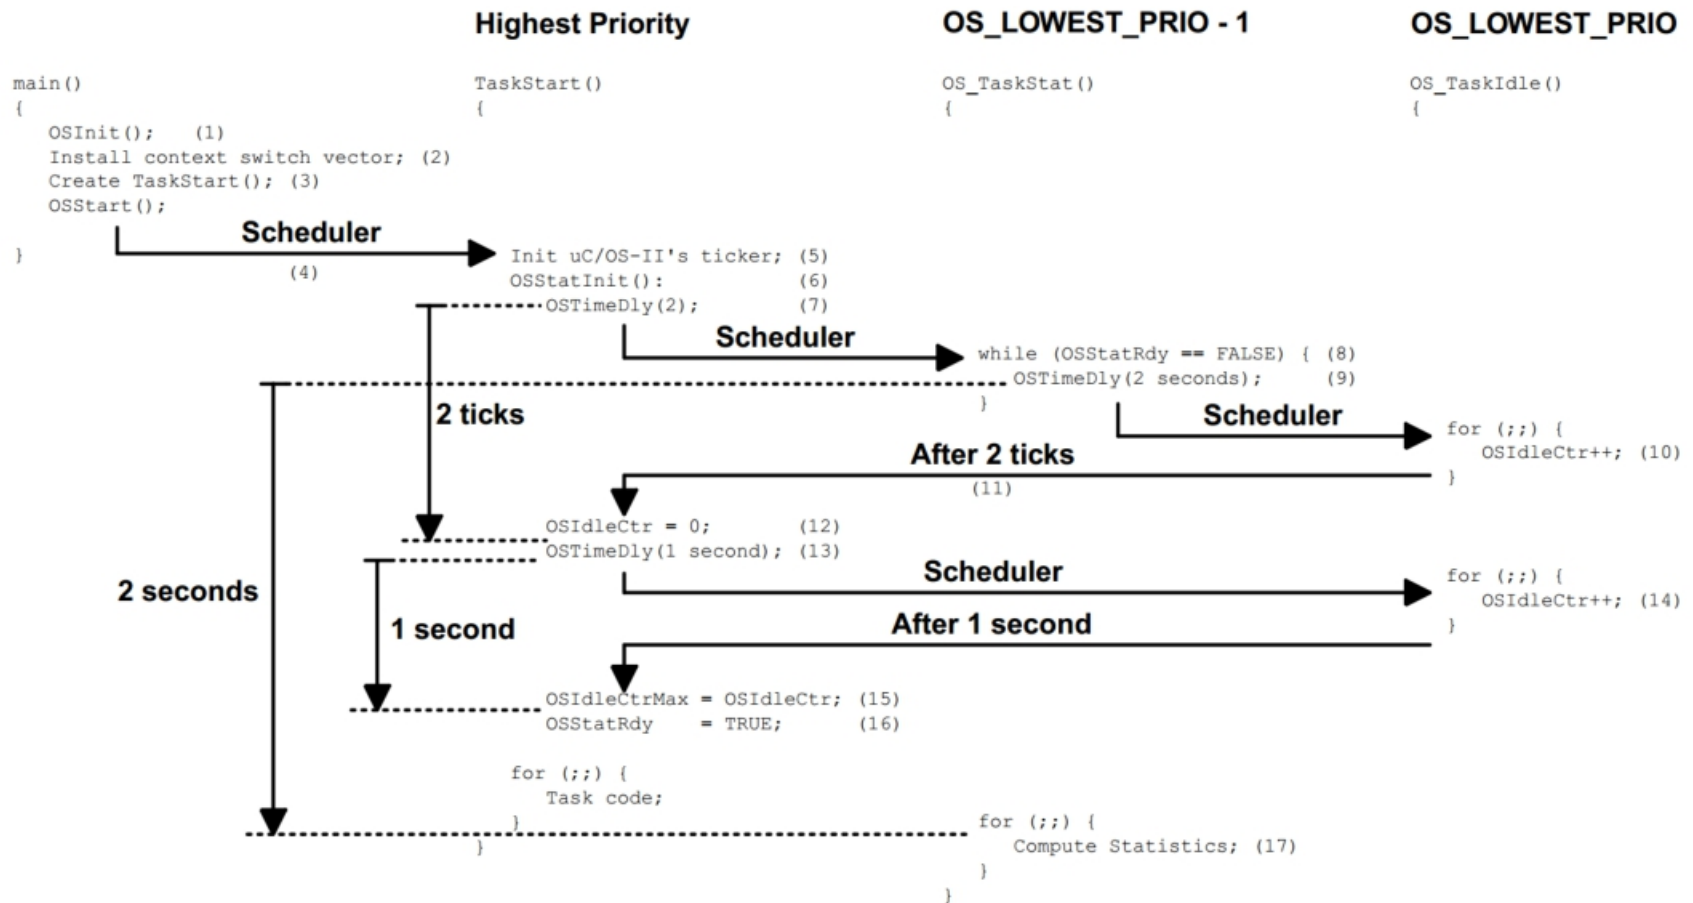
Statistics Task (cont'd)

- **To use OSCPUUsage, it must be initialised:**
 - Call OSStatInit() from *first and only* task running

```
void main(void){  
    // initialisation code here ..  
    OSTaskCreateExt(TaskStart, ...  
    OSStart();  
}  
void TaskStart(void *pdata){  
    OSStatInit();  
    // create other tasks and implement ...  
}
```

STATISTICS TASK INITIALISATION

```
void OSStatInit (void){  
    // must run when no other tasks running  
    // Finds OSIdleCtr for one second  
    OSTimeDly(2);  
        // now synchronised to timer interrupt  
    OS_ENTER_CRITICAL();  
        OSIdleCtr = 0L;  
    OS_EXIT_CRITICAL();  
    OSTimeDly(OS_TICKS_PER_SEC);  
    OS_ENTER_CRITICAL();  
        OSIdleCtrMax = OSIdleCtr;  
        OSStatRdy = TRUE;  
    OS_EXIT_CRITICAL();  
} // OSStatInit
```



How it works!!!

Study these two slides carefully

- F3.9(1) The first function that you must call in μ C/OS-II is `OSInit()`, which initializes μ C/OS-II.
- F3.9(2) Next, you need to install the interrupt vector that performs context switches. Note that on some processors (specifically the Motorola 68HC11), you do not need to install a vector because the vector is already resident in ROM.
- F3.9(3) You must create `TaskStart()` by calling either `OSTaskCreate()` or `OSTaskCreateExt()`.
- F3.9(4) After you are ready to multitask, call `OSStart()`, which schedules `TaskStart()` for execution because it has the highest priority.
- F3.9(5) `TaskStart()` is responsible for initializing and starting the ticker. You want to initialize the ticker in the first task to execute because you don't want to receive a tick interrupt until you are actually multitasking.
- F3.9(6) Next, `TaskStart()` calls `OSStatInit()`. `OSStatInit()` determines how high the idle counter (`OSIdleCtr`) can count if no other task in the application is executing. A Pentium II running at 333MHz increments this counter to a value of about 15,000,000. `OSIdleCtr` is still far from wrapping around the 4,294,967,296 limit of a 32-bit value. At the rate processor speeds are getting, it will not be too long before `OSIdleCtr` overflows. If overflow becomes a problem, you can always introduce some software delays in `OSTaskIdleHook()`. Because `OS_TaskIdle()` really doesn't execute any useful code, it's OK to throw away CPU cycles.
- F3.9(7) `OSStatInit()` starts off by calling `OSTimeDly()`, which puts `TaskStart()` to sleep for two ticks. This action is done to synchronize `OSStatInit()` with the ticker. μ C/OS-II then picks the next highest priority task that is ready to run, which happens to be `OS_TaskStat()`.
- F3.9(8) The code for `OS_TaskStat()` is discussed later, but as a preview, the very first thing `OS_TaskStat()` does is check to see if the flag `OSStatRdy` is set to `FALSE` and then delays for two seconds if it is.

How it works!!!

Study these two slides carefully

- F3.9(9) It so happens that `OSStatRdy` is initialized to `FALSE` by `OSInit()`, so `OS_TaskStat()` in fact puts itself to sleep for two seconds. This action causes a context switch to the only task that is ready to run, `OS_TaskIdle()`.
- F3.9(10) The CPU stays in `OS_TaskIdle()` until the two ticks of `TaskStart()` expire.
- F3.9(11)
- F3.9(12) After two ticks, `TaskStart()` resumes execution in `OSStatInit()`, and `OSIdleCtr` is cleared.
- F3.9(13) Then, `OSStatInit()` delays itself for one full second. Because no other task is ready to run, `OS_TaskIdle()` again gets control of the CPU.
- F3.9(14) During that time, `OSIdleCtr` is continuously incremented.
- F3.9(15) After one second, `TaskStart()` is resumed, still in `OSStatInit()`, and the value that `OSIdleCtr` reached during that one second is saved in `OSIdleCtrMax`.
- F3.9(16)
- F3.9(17) `OSStatInit()` sets `OSStatRdy` to `TRUE`, which allows `OS_TaskStat()` to perform a CPU usage computation after its delay of two seconds expires.

```
main()
{
  OSInit(); (1)
  Install context switch vector; (2)
  Create TaskStart(); (3)
  OSStart();
}
```

```
TaskStart()
{
  Init uC/OS-II's ticker; (5)
  OSStatInit(); (6)
  OSTimeDly(2); (7)
  while (OSStatRdy == FALSE) { (8)
    OSTimeDly(2 seconds); (9)
  }
  for (;;) {
    OSIdleCtr++; (10)
  }
  OSIdleCtr = 0; (12)
  OSTimeDly(1 second); (13)
  OSIdleCtrMax = OSIdleCtr; (15)
  OSStatRdy = TRUE; (16)
  for (;;) {
    Task code;
  }
  for (;;) {
    Compute Statistics; (17)
  }
}
```

```
OS_TaskStat()
{
  OS_TaskIdle()
}
```

```
OS_TaskIdle()
{
}
```

Highest Priority

OS_LOWEST_PRIO - 1

OS_LOWEST_PRIO

Scheduler

2 ticks

After 2 ticks

Scheduler

1 second

After 1 second

Scheduler

2 seconds

NB: If 100% CPU usage, this task may only run once at start up – can then report 0% CPU usage!

Computing CPU Usage

- When non idle tasks run, *IdleCtr* cannot increment.
- (*IdleCtrMax* – *IdleCtr*) used to calculate *CPUUsage* every second in *OS_TaskStat()* :

$$\begin{aligned} \text{OSCPUUsage\%} &= 100 \times \left(1 - \frac{\text{OSIdleCtr}}{\text{OSIdleCtrMax}} \right) \\ &= 100 - \frac{100 \times \text{OSIdleCtr}}{\text{OSIdleCtrMax}} \quad \leftarrow \text{Overflow possible} \\ &= 100 - \frac{\text{OSIdleCtr}}{\left(\frac{\text{OSIdleCtrMax}}{100} \right)} \quad \text{Integer divides now usable} \end{aligned}$$

void OS_TaskStat (void *p_arg)


```
void    OS_TaskStat (void *p_arg)
{
    while (OSStatRdy == OS_FALSE) {
        OSTimeDly(2 * OS_TICKS_PER_SEC);
    }
    OSIdleCtrMax /= 100L;
    for (;;) {
        OS_ENTER_CRITICAL();
        OSIdleCtrRun = OSIdleCtr;

        OSIdleCtr      = 0L;
        OS_EXIT_CRITICAL();
        OSCPUUsage      = (INT8U)(100L - OSIdleCtrRun /
OSIdleCtrMax);
        OSTaskStatHook();
        OSTimeDly(OS_TICKS_PER_SEC)
    }
}
```

Measuring Timer Interrupt Overhead

- OSIdleCtrMax

- Represents one second's worth of the idle task running *with interrupts occurring* .
- Eg with 2 interrupts per second:



Interrupt Overhead (cont'd)

- Suppose OSIdleCtrMax reaches **Idle2** counts in a second for 2 Hz interrupts
- Suppose **Idle1** counts for 1 Hz interrupts – measured after reconfiguring system.

Interrupt Overhead (cont'd)

- **Q: What is the %CPU overhead for the 1Hz timer interrupt?**
- **A: Idle1 – Idle2 represents one ISR time in idle count units.**

**=> Idle counts in an UNINTERRUPTED second
= Idle1 + (Idle1 - Idle2)**

%CPU time of One ISR per second is:

$$(Idle1 - Idle2) / [Idle1 + (Idle1 - Idle2)] * 100$$

- **Can you generalise this approach to 100 Hz interrupts and 1000 Hz interrupts?**